

sprintSix

December 9, 2024

1 Sprint 6: Building a Country Classifier Bot with ResNet

Owen Kroeger

- [Video Link](#) -

- [Jira Link](#) -

- [GitHub Link](#) - Using the ResNet architecture and a dataset of labeled geolocation images, we train, validate, and test a model while exploring its predictions.

Key Objectives: - Understand the dataset and preprocessing steps. - Train a ResNet-based classifier on geolocation data. - Evaluate the model and visualize its predictions.

1.1 Dataset Overview

We are using a dataset of ~50,000 Google Street View images organized into folders by country. The dataset was split into: - **Training set:** 70% of the images, used to train the model. - **Validation set:** 15% of the images, used to tune hyperparameters and evaluate the model during training. - **Test set:** 15% of the images, used to evaluate the final performance of the trained model.

The images were preprocessed to ensure consistent dimensions and normalization.

```
[ ]: # Importing libraries
import os
from torchvision import datasets, transforms

# Data transformations
transform = transforms.Compose([
    transforms.Resize((128, 128)), # Resize to consistent dimensions
    transforms.ToTensor(),         # Convert to Tensor
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
])

# Path to dataset
DATA_DIR = "./dataset/compressed_dataset"

# Load dataset
```

```
dataset = datasets.ImageFolder(DATA_DIR, transform=transform)
class_names = dataset.classes

print(f"Number of classes: {len(class_names)}")
print(f"Classes: {class_names[:10]}..." ) # Show a few class names
```

1.2 Model Architecture: ResNet

ResNet (Residual Network) is a widely used convolutional neural network (CNN) known for its ability to learn deep features effectively. We are using a pre-trained ResNet model (ResNet-18) and modifying its final layer to classify the images into countries.

```
[ ]: import torch
import torch.nn as nn
from torchvision.models import resnet18

# Initialize model
model = resnet18(weights="IMAGENET1K_V1")
model.fc = nn.Linear(model.fc.in_features, len(class_names)) # Adjust output_
                    ↪layer

# Display model architecture
print(model)
```

1.3 Training the Model

We train the ResNet model on the training dataset for a specified number of epochs. The training process involves: 1. **Forward Pass**: Feed images into the model to calculate predictions. 2. **Loss Calculation**: Compute the difference between predictions and actual labels using CrossEntropyLoss. 3. **Backward Pass**: Update model parameters to minimize the loss using the Adam optimizer.

We validate the model after each epoch to monitor its performance.

```
[ ]: # Import necessary libraries for training
import torch.optim as optim
from torch.utils.data import DataLoader, random_split
from tqdm import tqdm

# Split dataset
train_size = int(0.7 * len(dataset))
val_size = int(0.15 * len(dataset))
test_size = len(dataset) - train_size - val_size
train_dataset, val_dataset, test_dataset = random_split(dataset, [train_size,
                    ↪val_size, test_size])

# Data loaders
train_loader = DataLoader(train_dataset, batch_size=16, shuffle=True)
```

```

val_loader = DataLoader(val_dataset, batch_size=16)

# Loss function and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)

# Training loop
def train_one_epoch(epoch):
    model.train()
    running_loss = 0.0
    correct = 0
    total = 0
    for inputs, labels in tqdm(train_loader, desc=f"Epoch {epoch+1} Training"):
        inputs, labels = inputs.to("cpu"), labels.to("cpu") # Use "cpu" for
        ↪ this demonstration

        # Forward pass
        outputs = model(inputs)
        loss = criterion(outputs, labels)

        # Backward pass
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        # Metrics
        running_loss += loss.item() * inputs.size(0)
        _, preds = torch.max(outputs, 1)
        correct += (preds == labels).sum().item()
        total += labels.size(0)

    return running_loss / len(train_loader.dataset), 100 * correct /
    ↪ len(train_loader.dataset)

# Train for 1 epoch (for demonstration)
train_loss, train_acc = train_one_epoch(0)
print(f"Train Loss: {train_loss:.4f}, Train Accuracy: {train_acc:.2f}%")

```

1.4 Evaluating the Model

Once the model is trained, we evaluate it on the test dataset to measure its accuracy. We also visualize predictions on random test images.

```

[ ]: import matplotlib.pyplot as plt

# Evaluation on test data
def evaluate_and_visualize(model, test_loader, class_names):

```

```

model.eval()
images, labels = next(iter(test_loader))
with torch.no_grad():
    outputs = model(images)
    _, preds = torch.max(outputs, 1)

# Plot predictions
fig, axes = plt.subplots(1, 4, figsize=(15, 5))
for i, ax in enumerate(axes):
    image = images[i].permute(1, 2, 0).numpy() # Convert to HWC format
    image = (image * [0.229, 0.224, 0.225]) + [0.485, 0.456, 0.406] #
    ↪ Denormalize
    image = (image * 255).astype("uint8")
    ax.imshow(image)
    ax.axis("off")
    ax.set_title(f"True: {class_names[labels[i]]}\nPred:
    ↪ {class_names[preds[i]]}")
    plt.show()

# Run evaluation
test_loader = DataLoader(test_dataset, batch_size=4)
evaluate_and_visualize(model, test_loader, class_names)

```

1.5 Challenges

- **Class Imbalance:** The dataset contains more images for certain countries, which may lead to biased predictions.
- **Resource Constraints:** Training deep learning models on a CPU is significantly slower than using GPUs.

1.6 Next Steps

1. Experiment with data augmentation techniques to improve model generalization.
2. Implement techniques to handle class imbalance, such as oversampling or weighted loss.
3. Fine-tune the model on a smaller set of countries to reduce computational requirements.
4. Deploy the model and integrate it into the GeoGuessr bot for real-time predictions.

